




Learn & Ship Faster in An AI-Driven PDLC

BPM Dallas 
May 26th, 2026



BPM Intros




Host: A.J. Angus

Principal PM at Axon

Working on **Responsible AI** in Public Safety

Co-lead **BPM Dallas** Chapter



Speaker:
Buddy Triplett

Principal Product Manager at SpecBooks

Working on **Construction Software**

Leading **Lenny's Dallas** Meetups

Black Product Managers Network



Background

- Who we are: One of largest global networks of Product Managers with 2000+ members at all levels from aspiring PM to Chief Product Officer
- Mission: Support the upskilling and career development of Product Managers
- Vision: Product Managers have access to a network that champions their development

Three Focus Areas

- Skill Acquisition - Leveling up our members to earn more, and get promoted faster
- Community - Creating the closest, most sought after professional support network
- Growth Opportunities - Influencing the industry to support and hire more PMs

Join us by applying via <https://www.blackproductmanagers.com/apply>

VIRTUAL WORKSHOP

How I Claude

Learn & Ship Faster in
An AI-Driven PDLC

90 minutes · 7 modules · 7 hands-on artifacts



SCAN ME

One reframe. Better outcomes.

“PMs who ship faster don’t just write better prompts. They improve the system around the model — context files, validation loops, reusable workflows.”

The old pattern

PM reads inputs manually, synthesizes in their head, writes the brief, formats it three ways for three audiences.

It’s slow, inconsistent, doesn’t scale.

The new pattern

Claude is your connective tissue across the PDLC.

Not as one-off prompt receptacle, but as repeatable workflow engine from raw signal to shippable artifacts.

The foundation everything else runs on

These small, opinionated files beat one giant doc.

Load these once into a Claude Project. Every analysis, synthesis, and ticket from here on starts with the same strategic foundation — no re-explaining, no drift.

What makes context 'opinionated':

It includes what you're NOT doing. Stale or vague context degrades every output downstream.

DELIVERABLE

A loaded Claude Project with 3 context files

company.md

Mission, business model, strategic priorities, and explicitly what you're NOT doing this year

product.md

Core features, known pain points, current quarter focus, key tech constraints

personas.md

Named users, their JTBD, frustrations, and what success looks like for each one

Three inputs. Three different stories.

Each one is telling you something real. They don't agree. That's the point.

LEADERSHIP MEMO

Cross-project reporting dashboard

CEO requesting a dashboard showing drawing status, approvals, and overdue tasks across all projects — with CSV export.

1 account asking loudly

\$340K ARR renewal

45-day deadline

Rob Callahan, CEO · Meridian Group

41 SUPPORT TICKETS · 7 DAYS

Mobile drawing sync is broken

Field teams served stale or draft drawings on mobile. Wrong specs used on-site. One superintendent used a draft drawing for a concrete pour.

6 critical escalations

Active P1 incidents

\$1,200 material loss (1 site)

Field superintendents · Mid-market segment

POSTHOG ANALYTICS · 90 DAYS

Task dependency data is silently corrupting

Completing a parent task silently deletes dependent tasks. Most affected users don't know it's happening. Invisible unless you're watching the work flow funnels.

1,400–1,800 tasks lost silently

201 affected users

3 renewal accounts at risk

Active users · Most don't know yet

Synthesis ≠ summarization

✗ Weak prompt

“Summarize these support tickets for me.”

What you get:

- A paragraph-per-ticket summary
- No signal about what matters most
- No contradictions surfaced
- No confidence levels
- You still have to synthesize yourself



✓ Strong prompt

*“Identify 5–8 recurring themes. For each: ticket count, a representative quote, a confidence level, and any tickets that **contradict** this theme. Flag weak signals explicitly.”*

What you get:

- 5–8 named themes with counts
- Evidence quotes per theme
- Confidence levels on weak signals
- Contradictions flagged by name

The contradiction habit:

Ask Claude to name tickets that disagree with each theme. If support says X is broken but PostHog shows 2% hit rate — that tension is the insight.

DELIVERABLE

Themes table with evidence, confidence, and contradictions

One continuous Claude session

Rank → brief → constraint pass. No breaks, no switching tools. This is how fast shipping actually feels.

1 Quick rank

Drop themes into a simple impact/effort frame. Ask Claude to rank with reasoning. Takes 2 minutes.

Claude ranks. You override at least one. Write one sentence of rationale. That sentence is the PM judgment this module is training.

2 Stress-test

Ask Claude to argue against your top-ranked initiative. What evidence would flip this?

Prompt: "Argue against #1. What would have to be true for this to be the wrong bet?"

3 One-page brief

Problem statement · Target persona · Proposed approach · Key risks. All from the same session.

Link each section to evidence from the themes table. No unsupported claims.

4 Constraint pass

Inject a real constraint: "Eng says we can't touch auth this quarter." Ask Claude to redesign.

This step shows compounding value: Claude redesigns using the full context of the session — not just the last prompt.

DELIVERABLE

Ranked shortlist + solution brief v2 that survived one constraint pass

Prototypes are for demos. Flat screens are for engineers.

Empty State




What the user sees before data exists or on first load

Loading State




Skeleton layout with correct dimensions — no spinners in the void

Success State



Populated UI with realistic data and primary action visible

Error State



Failure messaging, clear recovery action, no dead ends

Why not a prototype?

Engineers can't pull interactive prototypes apart, can't reference a specific state in a ticket, and can't see edge cases without running the thing.

What Claude builds

A single HTML artifact: each screen as a labeled section with inline annotations. Static, referenceable, screenshot-able.

The ticket payoff

"Implement the empty state in Screen 2. CTA placement and copy must match the annotation." — That's the whole AC.

Brief → tickets in one prompt

Anatomy of a Claude-generated ticket

Context

Why this ticket exists — link to the brief and theme evidence

Acceptance Criteria

Specific, testable. Reference flat screen state by name.

Dependencies

What ships first or in parallel — be explicit, no assumptions

Edge Cases

States Claude flagged: empty, error, partial data, concurrent user

DELIVERABLE

3-5 implementation-ready tickets referencing flat screens

The unlock

Tickets that reference flat screens have shorter ACs. A screen says what copy can't.

"Implement the empty state in Screen 2. CTA, copy, and icon must match the annotation."

Scope check:

Ask Claude to split any ticket estimated over 3 days. Big tickets that sneak through here slow the sprint before it starts.

Make it run every month without starting from scratch

The shift:

You just ran this workflow manually. Now you'll write the recipe that makes Claude run it automatically — every month, with fresh data, no re-explaining.

Trigger file

Tells Claude what inputs to expect, what context files to load, and what confidence threshold flags for human review.

`run-synthesis.md`

Standing prompt

Module 2's synthesis prompt, templated with `{{DATE_RANGE}}`, `{{PRODUCT_AREA}}`, `{{QUARTER_FOCUS}}` placeholders.

`monthly-synthesis-prompt.md`

Output template

The format Claude always returns: themes table, top 3 signals, contradictions, recommended bet, PM action required.

`synthesis-output-template.md`

Evals check

Asks Claude to score its own output on 4 dimensions before you act on it. Catches drift before it reaches the roadmap.

`eval-rubric.md`

Evals habit:

After every synthesis run, ask Claude: "Score this output 1–5 on theme clarity, evidence quality, contradiction completeness, and recommendation confidence. Flag the two claims most likely to be wrong."

You need a way to grade Claude, not just use it.

Theme clarity

Are the themes specific enough to act on, or so broad they could mean anything?

1 = vague / 5 = actionable

Evidence quality

Is each theme backed by verbatim quotes and ticket IDs, or just Claude's summary?

1 = assertions / 5 = citations

Contradiction coverage

Did Claude find at least one signal that challenges each major theme?

1 = none found / 5 = all surfaced

Recommendation confidence

Is the 'recommended bet' traceable to specific evidence, or is it a guess?

1 = unsupported / 5 = fully evidenced

If any dimension scores below 3:

Don't act on the output yet. Identify which prompt produced the weak output and tighten it before the next run.

Everything you built today is transferable.

These aren't personal productivity tools. They're team infrastructure.

What to share and where

 Context pack Claude Project shared with team	Every PM and designer who touches this product
 Prompt library Notion / Confluence / shared doc	Anyone running synthesis, writing tickets, or drafting briefs
 Automation recipe GitHub repo or shared drive	The PM who owns the monthly synthesis rotation
 Eval rubric Pinned in your team's Claude Project	Anyone reviewing Claude outputs before acting on them

One ask: before you leave today, identify one teammate who would benefit from the prompt library. Send it to them this week.

You built 7 artifacts today.

And a workflow that runs again next month — by you or anyone on your team.

01

Context Pack

3 opinionated files in a Claude Project

02

Themes Table

Evidence · confidence · contradictions

03

Ranked Shortlist

PM rationale on the top bet

04

Solution Brief v2

Survived one constraint pass

05

Flat Screen Artifact

4 annotated states for engineering

06

Ticket Set

3–5 implementation-ready tickets

07

Automation Recipe

Trigger · prompt · template · eval rubric

Take-homes:

Context file templates · Full prompt library · Flat screens prompt · Eval rubric

Evals

Grade outputs before acting.

Team

Share context pack + prompts.

Repeat

Real data run within 2 weeks.

Lenny's Newsletter Meetup





You have manage access for this event. [Manage](#)

 Presented by **Lenny's Newsletter ...** [Subscribe](#)

All meetups are hosted by members of the Friends of Lenny's Newsletter community! Learn more at lennysnewsletter.com/community 🔥

Lenny's Newsletter Meetup in Dallas (with free drinks courtesy of Stella)

Wednesday, May 27
5:30 PM - 8:00 PM

The Rackhouse
The Colony, TX

[Add to Calendar](#) [Share](#) [More](#)

You're In Starting in 1d

[No longer able to attend? Notify the host by cancelling your registration.](#)

Get Ready for the Event
Profile Complete - Reminder: Email

Thank you!

Let's Connect on LinkedIn



SCAN ME

Appendix

Advanced tips for AI-native PMs

Three techniques for PMs who already have a workflow (plus a bonus)
and want to make it significantly more powerful.

01

Teach Claude your taste

02

Audit before you read

03

Chain, don't restart

Teach Claude your taste, not just your task.

The problem

Most PMs use Claude transactionally — describe the task, get an output, edit it to sound like themselves. You're re-teaching your voice from scratch every single conversation.

The technique

Paste 3–4 pieces of writing you're proud of. Ask Claude to identify your patterns and write a style guide you can reuse in CONTEXT.md.

```
Analyze these examples of my writing.  
Identify the specific patterns –  
sentence structure, level of detail,  
what I include vs. omit, how I handle  
uncertainty, how I open and close.
```

```
Write a style guide I can paste into  
future conversations to reproduce  
this voice.
```

The delta between "Claude's version" and "my version" often collapses. You edit for substance, not style.

Use Claude to audit its own output before you do.

The problem

Most people treat Claude's first output as the thing to edit instead of reading the whole document looking for problems. That's the slow, unfocused way.

The technique

Send a second message immediately after generation, before you read the output carefully. Ask Claude to surface its own failure modes first.

```
Before I review this – audit it yourself.
```

- What did you have to assume or invent because I didn't give you enough info?
- Where are you least confident in the accuracy or specificity?
- What would a skeptical engineer / VP / customer push back on first?
- What's the weakest sentence?

You're not reading the whole doc looking for problems. You go straight to the flagged sections. Works on anything.

Chain outputs, don't restart conversations.

The problem

Intermediate users start a new Claude conversation for every task. No shared context. Every output is an island. The exec brief can contradict the PRD.

The technique

Explicitly hand each output forward as the authoritative source for the next task in the same conversation (context window). Build a pipeline, not a pile of tabs.

```
Use this memo as the authoritative
source for everything that follows.
Do not introduce information not
present in it.
```

```
First task: write the exec brief.
After I approve it → eng handoff.
After that → Slack update.
```

One session. Exec brief, eng handoff, Slack update — all consistent with each other. No more "the brief says X but the PRD says Y."

Bonus concept

The attention budget.

Straight from Anthropic's engineering team.

"Context must be treated as a finite resource with diminishing marginal returns. Every new token introduced depletes the [attention] budget by some amount, increasing the need to carefully curate the tokens available to the LLM."

— Anthropic Engineering, *Effective Context Engineering for AI Agents* (2025)

What this means for PMs

Too little context

Claude guesses. Generic output. Could apply to any company, any product.

Too much context

Attention diluted. The specific thing you needed gets lost in the noise.

Goldilocks

Smallest set of high-signal context that fully defines the task.

The test: would a smart new hire need this to do the task? If no — leave it out.